

[prints](#)[Fellowships](#)[Research](#)[Impa](#)[Donate →](#)[Get Started →](#)[JUN 19, 2026 - JUN 21, 2026](#)[ONLINE & IN-PERSON](#)

Global South AI Safety Hackathon

Bringing the Global South into AI safety
June 19 - 21, 2026



\$6,000 in cash prizes +
Apart Fellowship pathway



Apart Research

x



Schmidt Sciences

AI safety research is concentrated in a handful of countries. This hackathon changes that. Build AI safety tools, evaluations, and policy research from Latin America, Africa, or Asia, compete within your region, and join a pipeline from hackathon to fellowship to placement. With Support from Schmidt Sciences.

This event is ongoing.

[Sign Up →](#)[Submit Your Project →](#)

1,479

Sign Ups

[Overview](#)[Resources](#)[Guidelines](#)[Schedule](#)

Guidelines



Judging Criteria

Dimension 1: Impact Potential & Innovation

How much would this matter for AI safety if it worked? How innovative is it?

For scores of 4-5: is this actually new to the field, or replicating recent work?

Score	Description
1	Negligible. No clear problem addressed, or no meaningful novelty.
2	Limited. Addresses a real

	problem but with a generic or well-trodden approach. Incremental at best.
3	Moderate. Clear problem with a reasonable approach; some novelty in framing or method beyond routine application of existing tools.
4	Significant. Important problem with an original approach, or identifies a neglected problem area. A valuable contribution others could build on.
5	Exceptional. Tackles a critical AI safety problem with a genuinely novel approach, or opens a new research direction. Clear theory of change.

You'd be excited to share this with researchers in the area.

Dimension 2: Execution Quality

How sound are methodology, implementation, and findings?

Score	Description
1	Seriously flawed. Methodology broken, results uninterpretable, or implementation doesn't work.
2	Weak. Approach has significant gaps: missing validation, flawed experimental design, or incomplete implementation.
3	Competent. Technically solid given the short duration. Methodology makes sense,

	results are interpretable, limitations acknowledged, work builds toward clear conclusions.
4	Strong. Thorough methodology with convincing validation. Results clearly support conclusions. Immediately useful for future work.
5	Exceptional. Ambitious scope executed rigorously. Surprising findings, novel methods, or unusually robust validation.

Dimension 3: Presentation & Clarity

How clearly are work, findings, and impact potential communicated?

Score	Description
1	Incomprehensible. Cannot determine what the project is actually claiming or doing.
2	Hard to follow. Key information buried, missing, or diluted by excessive length. Significant effort to extract main points.
3	Clear enough. Can understand the problem, approach, and results without undue effort. Core content clearly present: problem, method, findings, limitations.
4	Well presented. Easy to follow, well-structured, appropriate level of detail. Target audience would get it quickly.
5	Exceptionally clear. A pleasure

to read. Complex ideas made accessible. Could serve as a model for how to present this type of work.

Submission Requirements

A complete submission includes:

1. A research report in PDF format using the official [submission template](#)
2. A project title and brief abstract (150 words max)
3. Author names and affiliations for all team members
4. An indication of which regional track and sub-track your project addresses

Important: Include a section called "Limitations and Dual-Use Considerations" that addresses limitations of your approach, potential misuse risks, and suggestions for future improvements.

Report structure (recommended):

There is no hard page limit. Most winning projects are 4-8 pages.

The report should include:

- Introduction: What problem did you address? Why does it matter for AI safety in your region or globally?
- Related Work: What existing work does your project build on?
- Methodology: What did you build or test? Describe your approach in enough detail for someone to replicate it
- Results: What did you find? Include quantitative results where possible
- Discussion: What are the implications? What are the limitations? What would you do with more time?
- References: Cite relevant prior work
- Do not submit the whole report in *Italic*

Important notes:

- You can submit as an individual or as a team
- You can build on existing work, but you **MUST** clearly identify what is **NEW** work done during the hackathon

- Submit through the official submission form on the hackathon page
- If you run into submission issues, DM Kamil on Discord or email sprints@apartresearch.com
- Have team names, emails, Discord handles, project title, abstract, and PDF ready before starting the form.
- Reports can be in English, Spanish, or Portuguese. English is preferred if your team is comfortable with it, but we will not penalize non-English submissions.
- Due to the high volume of submissions, we may not be able to provide written feedback for every project, but all projects will be evaluated.

**Submission
Template => [Link](#)**

**Discord Server Invite
=> [Link](#)**

**Frequently Asked
Questions**

Eligibility and Format

Q: Do I have to be from the Global South to participate?

We use Britannica's definition of the Global South. If you're based in Latin America, Africa, the Middle East, or Asia/Oceania (excluding Japan, South Korea, Taiwan, Australia, New Zealand), you're in scope.

If you're based outside these regions, you're still welcome: submit your project to the Open Track (Track 4). If you're not sure where you fit, sign up and we'll sort it out.

Q: Do I need a research background or coding experience to participate?

No formal requirement. Past sprints have had grad students, undergrads, industry engineers, and self-taught researchers all submit winning projects. What helps: ability to scope a small experiment, write up findings clearly in a short PDF. We provide resources, mentor support during the weekend, and the rubric publicly so you know what's being evaluated. If you're unsure whether you'd fit, ask in the Frequently Asked Questions

Answers below are from the Apart master FAQ database (Kamil-approved). For anything not covered, ask on Discord help-

desk or email
sprints@apartresearch.com.

Eligibility and Format

Q: Do I have to be from the Global South to participate?

We use Britannica's definition of the Global South. If you're based in Latin America, Africa, the Middle East, or Asia/Oceania (excluding Japan, South Korea, Taiwan, Australia, New Zealand), you're in scope. If you're not sure, sign up and we'll sort it out.

Q: Do I need a research background or coding experience to participate?

No formal requirement. Past sprints have had grad students, undergrads, industry engineers, and self-taught researchers all submit winning projects. What helps: ability to scope a small experiment, write up findings clearly in a short PDF. We provide resources, mentor support during the weekend, and the rubric publicly so you know what's being evaluated. If you're unsure whether you'd fit, ask in the Discord help-desk.

Q: Is the hackathon remote or in-person? Can I attend remotely if there's a local hub in my city?

Default is remote. We also have

local in-person hubs where people optionally gather for the weekend. You can join remotely from anywhere even if there's a hub in your city. Hubs are optional and additive, not a separate event. Talks are livestreamed to remote participants; judging is always asynchronous and online.

Q: When is the hackathon?

June 19-21, 2026. It starts Friday evening. Submissions close Sunday June 21 at 11:59 PM AoE (Anywhere on Earth).

Teams and Tracks

Q: Can I participate solo, or is it teams only? What's the max team size?

Both solo and team submissions are welcome. Maximum team size is 5. Smaller teams (1-3) are common and have won prizes. If you're solo and want a team, post in the Discord #team-formation channel a week or two before the event. We don't auto-match teams; it's organic.

Q: How strict are you on the tracks? Can my project fit between two or none of them?

Tracks are guidance, not gates. If your project genuinely lives between two tracks, pick the closest one and explain the cross-

track angle in your submission. Judges score on Impact, Execution, and Presentation across all tracks; nobody gets disqualified for fit. If it really fits none of the tracks but is in scope for the hackathon's overall topic, submit anyway and Kamil will route it. Better to submit and have us re-bin than to skip.

Q: Can a team have members from different regions?

Yes. Choose the regional track that best fits your project. Your team can include members from any location. You compete for prizes in whichever regional track you select.

Submissions

Q: What do I submit?

A research report in PDF format using the submission template. Think of it as a mini research paper documenting your problem, approach, results, and implications. Not a product demo.

Q: Can I build on existing work, or does everything have to be from scratch during the hackathon?

Yes, you can build on existing work, as long as the contribution you make during the hackathon is clearly identifiable in your submission. Judges look for what's

new this round. Cite prior work in your PDF (anywhere from a sentence to a methods paragraph; whatever makes the delta clear). The Project Eligibility note in our rubric covers this. Submissions that don't show a clear during-event delta tend to score lower on Execution.

Q: Do you recommend a particular referencing/citation style for the project write-up?

No required style. Pick whichever you are comfortable with (APA, MLA, Chicago, IEEE, etc.) and stay consistent throughout the paper. The judges care about the substance of the work, not the citation format.

Q: Can I submit an unfinished project?

Yes. Submitting something unfinished is always better than not submitting. Judges evaluate what you accomplished during the hackathon timeframe.

Judging, Results, and Prizes

Q: When are winners announced after the hackathon ends?

Roughly 1-3 weeks after the event. The timeline: judges receive assignments Monday or Tuesday after the weekend, reviews are due the following Sunday, then we

tally and rank, then we send winner notifications. For biosecurity-related events we add a hazard review step before announcing publicly, which can extend the timeline. We send everyone reviewer feedback after winners are notified.

Q: How are prizes awarded?

Prizes are awarded per region, not globally. Latin America has 3 winning teams, Asia has 2, and Africa has 1. You compete with participants in your regional track.

Q: How and when do prize payments arrive after the hackathon?

Winners are announced 1-3 weeks after the event. After announcement, your team agrees on a prize split and nominates one person to receive the funds. That nominee fills out a payment form and is onboarded to our payment processor (Ashgro/Ramp). Payment lands roughly 2-4 weeks after the form is submitted, depending on the payout method. The nominee handles distribution to the rest of the team. International transfers may take longer.

Support

Q: How do I get help during the hackathon?

Three ways:

- Discord help-desk channel (tag @Kamil Alaa)
- DM Kamil on Discord
- Email:
sprints@apartresearch.com

